

Language Understanding Systems

Language Modeling with OpenGRM Tools

Evgeny A. Stepanov

SISL, DISI, UniTN & VUI, Inc.
`evgeny.stepanov@unitn.it`

Outline

- 1 Language Models
- 2 OpenGRM for Language Modeling
- 3 Exercises


Section 1

Language Models

LM Functionality


LM Functionality

- Compute n-gram counts from the corpus



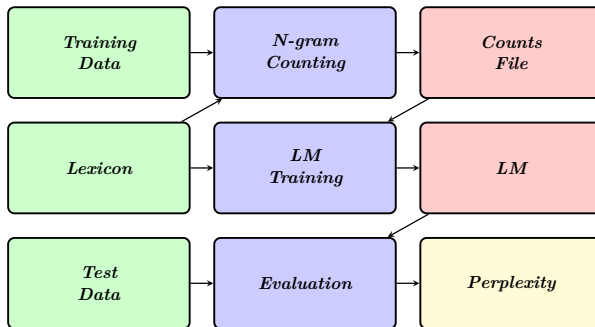
LM Functionality

- Compute n-gram counts from the corpus
- Train LM from the n-gram counts file

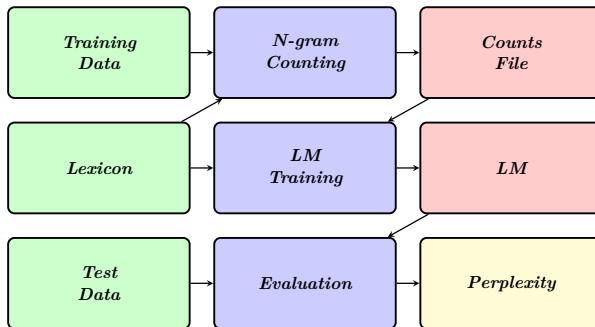


LM Functionality


- Compute n-gram counts from the corpus
- Train LM from the n-gram counts file
- Calculate the test data *perplexity* using LM



Language Modeling



LM Tools: OpenGRM



OpenGRM: Discounting

OpenGRM

- Absolute (Ney)
- Witten-Bell (default)
- Kneser-Ney
- Katz
- presmoothed
- unsmoothed

Section 2

OpenGRM for Language Modeling

OpenGRM Tools

OpenGRM Tools

- **ngramcount**: takes as input a *fst archive* representing a text input and outputs an *fst* representing counts of n-grams

OpenGRM Tools

- **ngramcount**: takes as input a *fst archive* representing a text input and outputs an *fst* representing counts of n-grams
- **ngrammake**: takes as input a *counts fst* (the output of ngramcount) and outputs an *fst* representing a n-gram back-off stochastic LM

OpenGRM Tools

- **ngramcount**: takes as input a *fst archive* representing a text input and outputs an *fst* representing counts of n-grams
- **ngrammake**: takes as input a *counts fst* (the output of ngramcount) and outputs an *fst* representing a n-gram back-off stochastic LM
- **ngrammerge**: merges two *counts fsts* (the output of ngramcount) or two *LMs* (the output of ngrammake)

OpenGRM Tools

- **ngramcount**: takes as input a *fst archive* representing a text input and outputs an *fst* representing counts of n-grams
- **ngrammake**: takes as input a *counts fst* (the output of ngramcount) and outputs an *fst* representing a n-gram back-off stochastic LM
- **ngrammerge**: merges two *counts fsts* (the output of ngramcount) or two *LMs* (the output of ngrammake)
- **ngramapply**: intersects n-gram model with *fst archive*;

OpenGRM Tools

- **ngramcount**: takes as input a *fst archive* representing a text input and outputs an *fst* representing counts of n-grams
- **ngrammake**: takes as input a *counts fst* (the output of **ngramcount**) and outputs an *fst* representing a n-gram back-off stochastic LM
- **ngrammerge**: merges two *counts fst*s (the output of **ngramcount**) or two *LMs* (the output of **ngrammake**)
- **ngramapply**: intersects n-gram model with *fst archive*;

Other useful tools:

- **ngramsymbols**: produces a lexicon from an input text corpus;
- **ngramread/ngramprint**: utilities to read/write n-gram models from/to text file;
- others...

FAR Tools

FARs – *weighted finite-state machines archives* – a concatenation of the file representation of one or more finite-state machines.

Commands:

- **farcompilestrings** – text to FAR
- **farprintstrings** – FAR to txt
- **farinfo** – information on contained FSMs
- **farequal** – compare two FARs for equality
- **farcreate** – create FAR from input FSTs
- **farextract** – split FAR into constituent FSTs

Refer to manual for options:

<http://www.openfst.org/twiki/bin/view/FST/FstExtensions#FstArchives>

Example

Input file: text.txt

Build Language Model:

```
ngramsymbols < text.txt > lex.txt
farcompilestrings --symbols=lex.txt
                  -keep_symbols=1 text.txt > text.far
ngramcount --order=3 text.far > text.cnts
ngrammake --method=witten_bell text.cnts > text.lm
```

Generate random sentence:

```
ngramrandgen text.lm | farprintstrings
```

Section 3

Exercises

Exercises

- Read Tool Manuals (for options)
- Train different language models (LM) using train.txt
 - vary order {1-3}
 - vary smoothing
 - take care of unknown words using lexicon (e.g. frequency cut-off)
 - compute LM perplexity on test.txt
 - report order & smoothing method with lowest perplexity